



kubernetes

Why kubernetes ?



Container Clusters

- What if we have 10s, 100s, 1000s of running containers on multiple VMs?
- How to deploy, scale, restart, manage all of these containers?
- What problems do they solve?
 - Management
 - Metrics
 - Health checks
 - Security
 - Abstraction of hardware
 - Networking
 - Scheduling
 - Scaling
 - Deployment
 - Rollbacks
 - Zero-downtime / blue-green
 - Service discovery

A Brief Kubernetes History

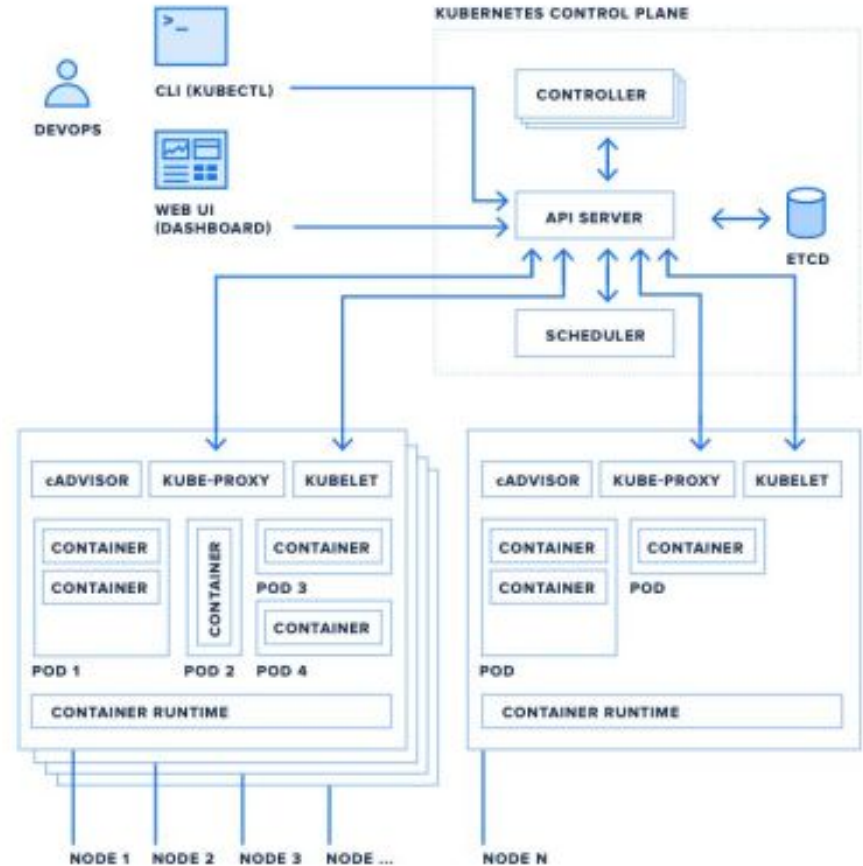
- “K8s”
- Evolved out of Borg (Google’s internal container cluster)
- Open sourced ~2014
- Grew in popularity, open source velocity increased
- Now the most popular container cluster (most cloud platforms have some sort of managed K8s offering)
- Features added regularly and frequently
- Cloud Native / CNCF - Kubernetes, Prometheus, Fluentd

Kubernetes Architecture

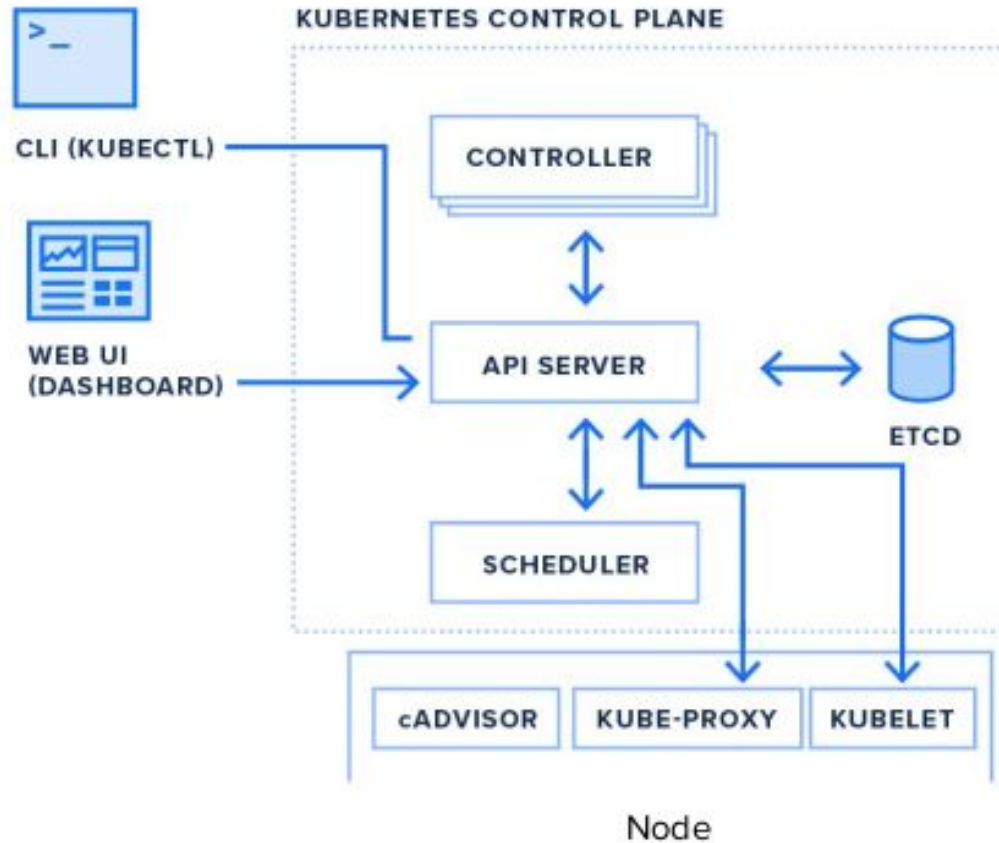
- Client side :- CLI(KUBECTL)

WEB UI(DASHBOARD)

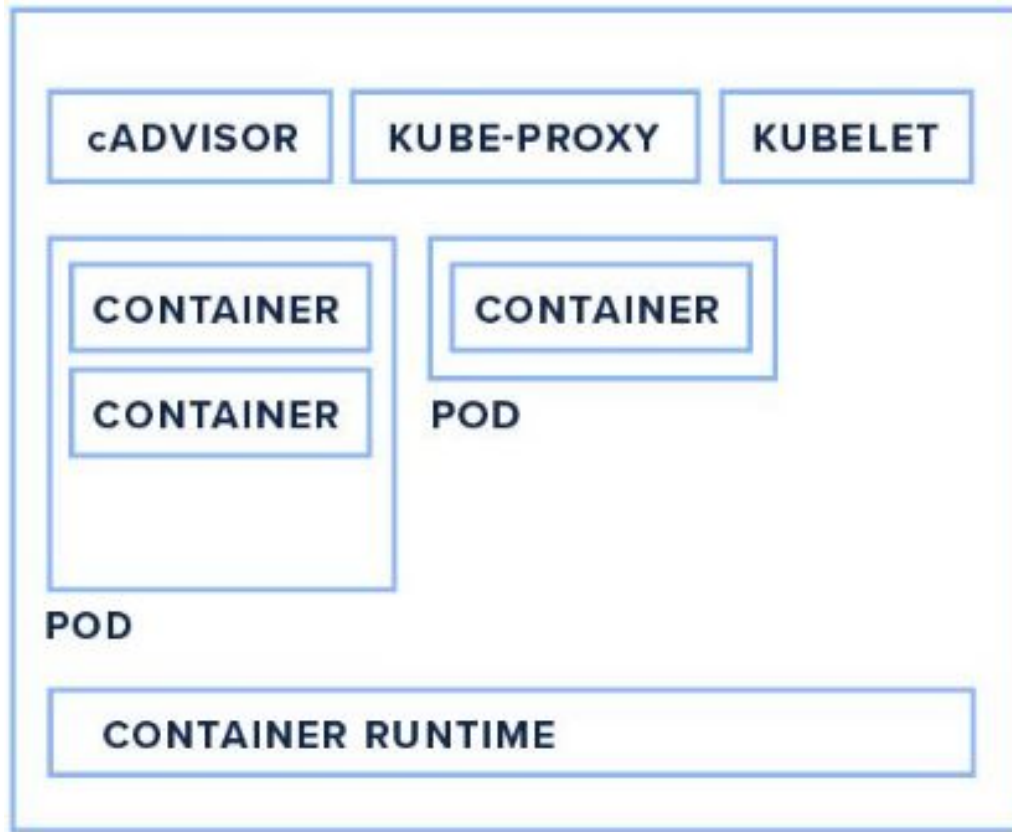
- Kubernetes Control plane or master node
- Minions node or worker node



Kubernetes control plane



Kubernetes worker nodes



Kubernetes installation

Single Node

- Docker desktop
- Minikube

Custom kubernetes

- Kubeadm
- Kubespray

Cloud

- AWS - EKS
- Azure - AKS
- Google - GKE

Minikube installation

What you'll need

- 2 CPUs or more
- 2 GB of free memory
- 20 GB of free disk space
- Internet connection
- Container or virtual machine manager, such as: [Docker](#), [QEMU](#), [Hyperkit](#), [Hyper-V](#), [KVM](#), [Parallels](#), [Podman](#), [VirtualBox](#), or [VMware Fusion/Workstation](#)

FOR linux

```
curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
```

```
sudo install minikube-linux-amd64 /usr/local/bin/minikube
```

For windows

<https://storage.googleapis.com/minikube/releases/latest/minikube-installer.exe>

If using powershell

```
New-Item -Path 'c:\' -Name 'minikube' -ItemType Directory -Force
```

```
Invoke-WebRequest -OutFile 'c:\minikube\minikube.exe' -Uri  
'https://github.com/kubernetes/minikube/releases/latest/download/minikube-windows-amd64.exe'  
-UseBasicParsing
```

FOR mac

curl -LO

<https://storage.googleapis.com/minikube/releases/latest/minikube-darwin-amd64>

sudo install minikube-darwin-amd64 /usr/local/bin/minikube

Some K8s commands

Minikube start

Minikube stop

Kubectl version

Kubectl get

Kubectl apply


Kubectl create

Kubectl delete

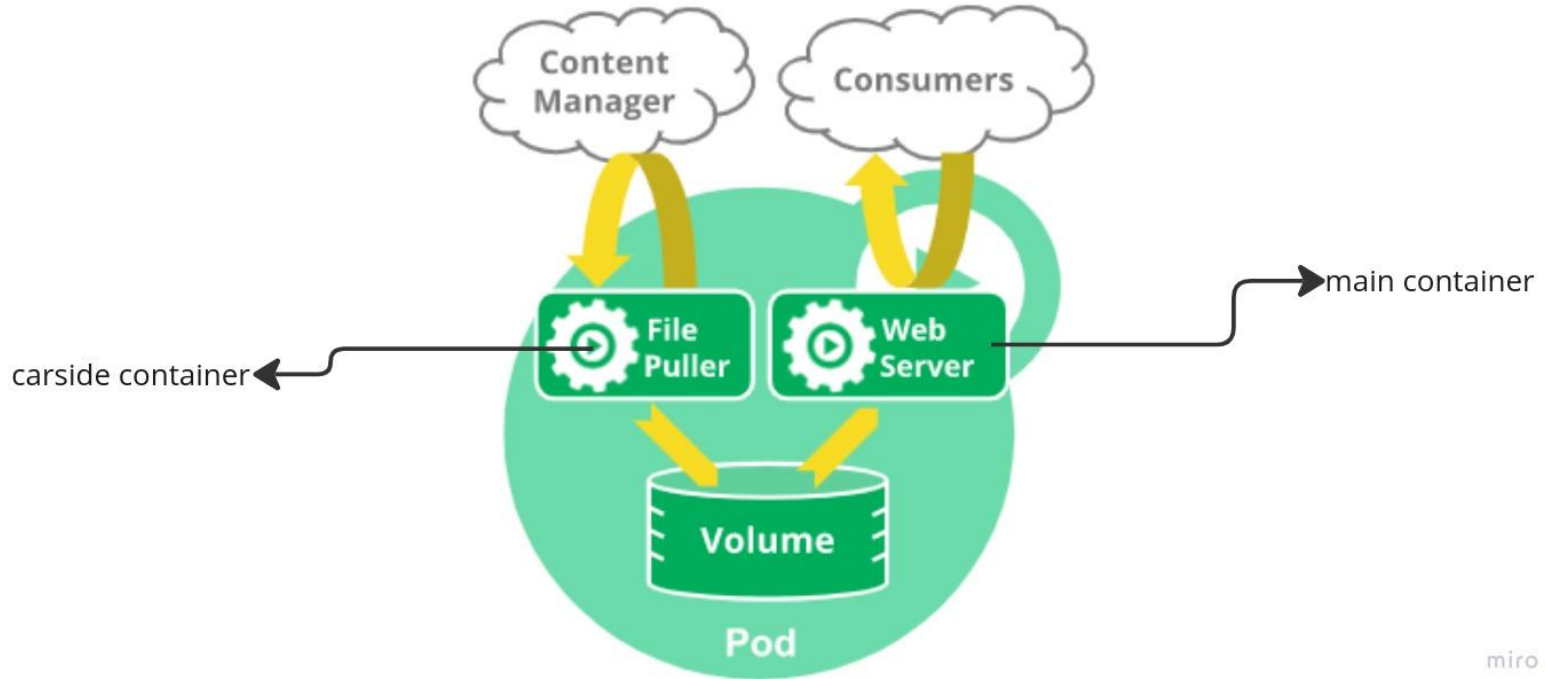
K8s components

PODS

- **Pods** are the smallest deployable units of computing that you can create and manage in Kubernetes.
- A **Pod** is a group of one or more containers.
- **Pods** that run a single container.
- **Pods** that run multiple containers that need to work together.
- **Pod** containers share resources
 - Storage
 - Network (localhost)
 - Always run on the same Node



Multiple containers in single pod



Create a container

Create a yaml file eg :- ak.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: akpod1 # name of your any kind type
spec: # to create env
  containers:
  - name: akc1
    image: nginx
    ports:
    - containerPort: 80
```

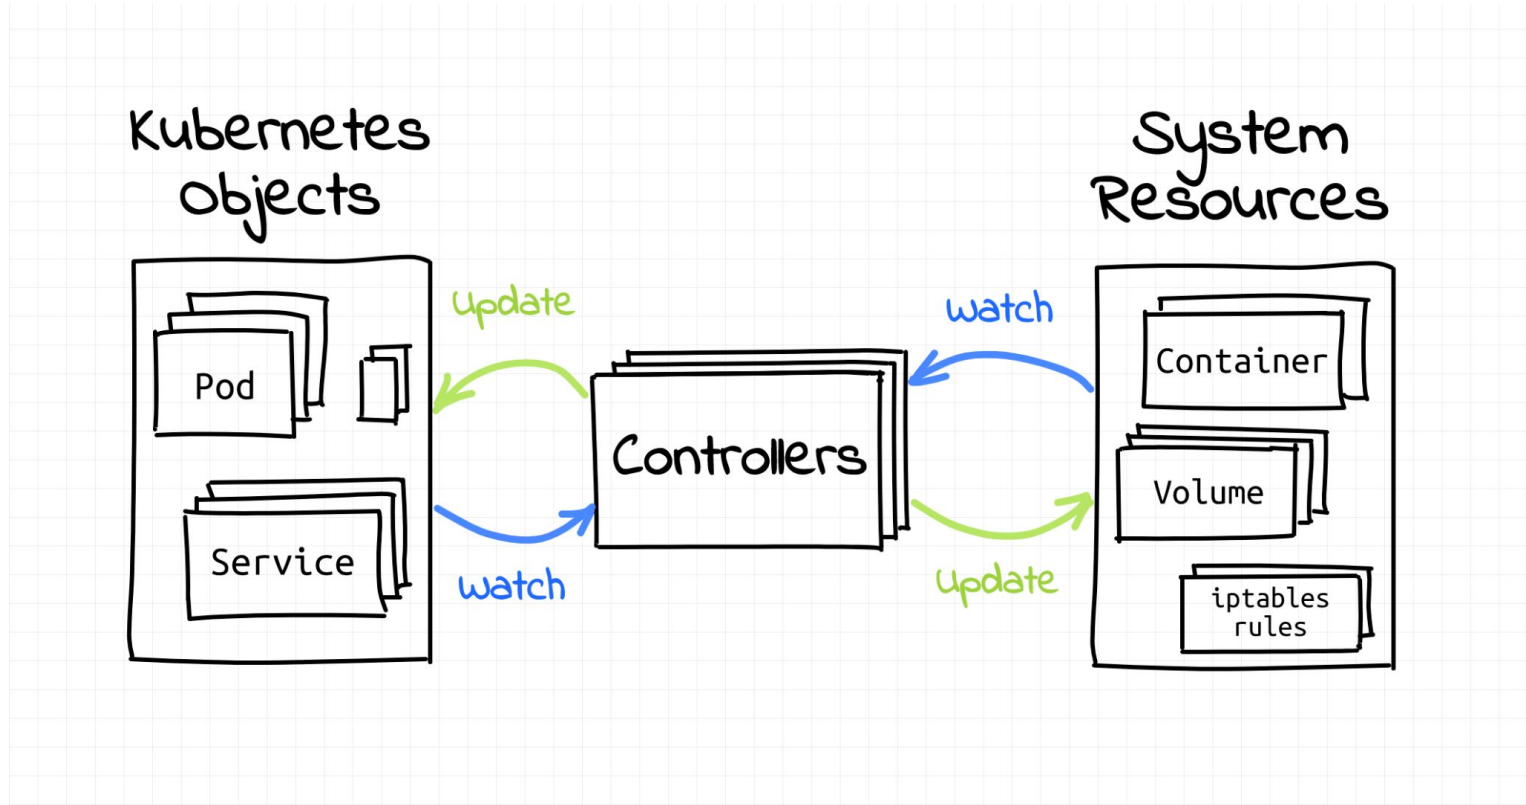


POD problems

- 1) Recrete [auto]
- 2) Scale(pod)
- 3) perhaps several Pods, to carry out a task and then stop.



Controller



K8s native Controller

Replication controller[RC]:-A ReplicationController ensures that a specified number of pod replicas are running at any one time. In other words, a ReplicationController makes sure that a pod or a homogeneous set of pods is always up and available.

ReplicaSet[RS]:- A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time.

ReplicationController

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: ashu-rc1
spec:
  replicas: 1 # number of pods
  template: # pod yaml info
    metadata:
      labels:
        x1: akash
    spec: # to create env
      containers:
        - name: ashuc1
          image: nginx
          ports:
            - containerPort: 80
```

Kubernetes workload according to apps

For Stateless app:- (eg :- Webapp)

- Deployments
 - ReplicaSets
 - Pods
 - Container

For stateful app:- (eg : Databases)

- StatefulSets

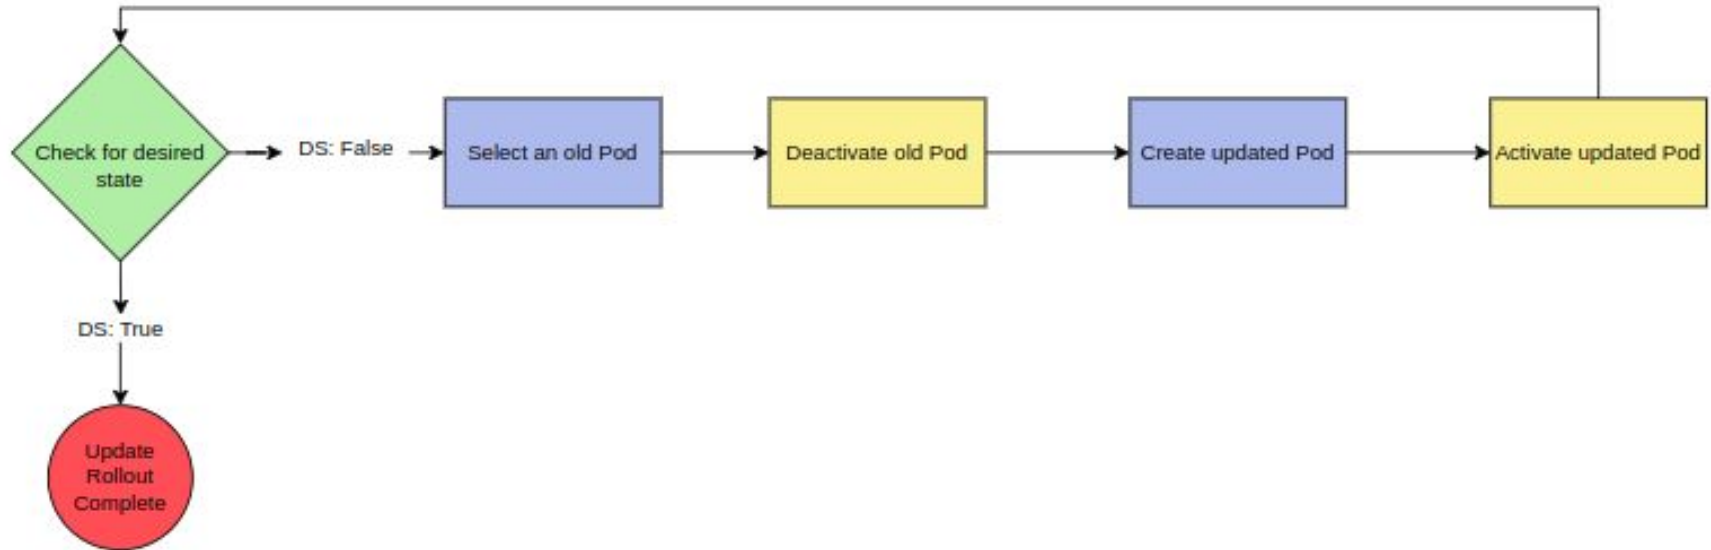
Deployments

A Deployment provides declarative updates for Pods and ReplicaSets.


Update Deployment Strategies

- **Rolling update strategy:** Minimizes downtime at the cost of update speed.
- **Recreation Strategy:** Causes downtime but updates quickly.
- **Canary Strategy:** Quickly updates for a select few users with a full rollout later.

Rolling update strategy

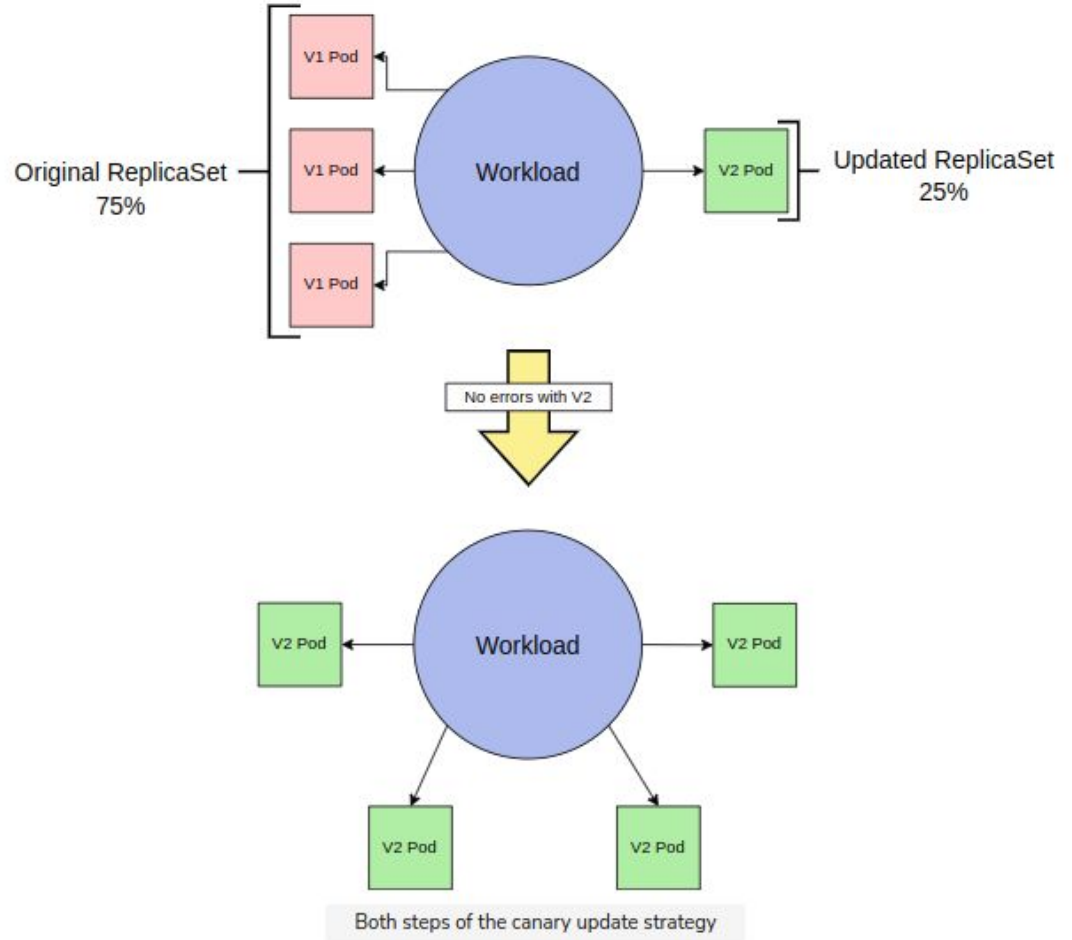


Recreation Strategy



Recreate update strategy flowchart

Canary Strategy



Deployments


```
kubectl create deployment  
akdep1 --image=nginx --port 80  
--dry-run=client -o yaml  
>deployment.yaml
```

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  creationTimestamp: null  
  labels: # label of deployment  
    app: akdep1  
  name: akdep1 # name  
spec:  
  replicas: 1 # number of pod  
  selector: #  
    matchLabels:  
      app: akdep1  
  strategy: {} # app upgrade strategy -- rolling updates  
  template: # to create pods  
    metadata:  
      creationTimestamp: null  
      labels: # label of pods  
        app: akdep1  
    spec:  
      containers:  
        - image: nginx  
          name: nginx  
          ports:  
            - containerPort: 80  
          resources: {}  
      status: {}
```

Kubernetes Networking



CNM(Container Network Model)



Docker



CNI

CNI(Container Network Interface)




- calico
- flannel
- AWS CNI
- Weave
- Romana
- ACI (cisco)
- Multos



CNM (container network model)

Company - docker

Runtime engine - docker



CNI(container networking model interface)

Company - CoreOS



Runtime engine - RKT



Service in k8s

Service is a method for exposing a network application that is running as one or more Pods in your cluster.

Each Service object defines a logical set of endpoints (usually these endpoints are Pods) .

Cluster IP


NodePort

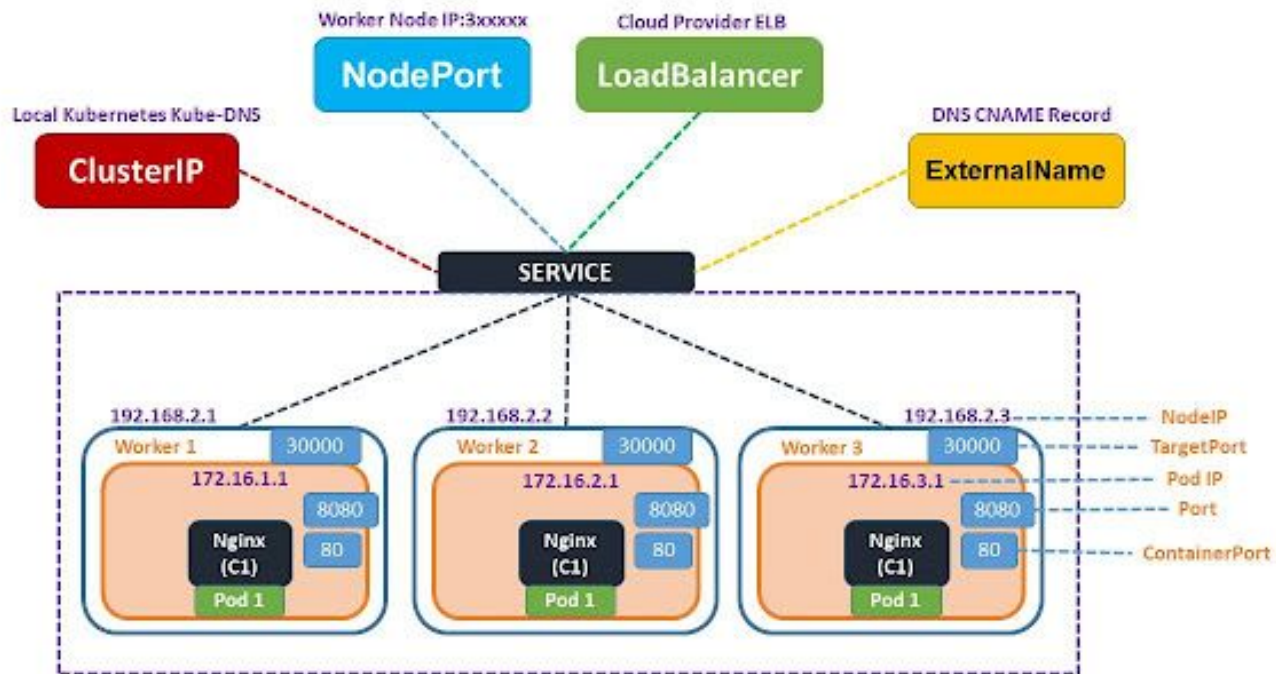
Kubernetes Services:
4 Types

LoadBalancer

ExternalName

NodePort





K8s secrets

A Secret is an object that contains a small amount of sensitive data such as a password, a token, or a key. Such information might otherwise be put in a Pod specification or in a container image.